

How to get peak FLOPS (CPU)
— What I wish I knew when I was twenty
about CPU —

Kenjiro Taura

Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

What you need to know to get a nearly peak FLOPS

- so you now know how to use multicores and SIMD instructions
- they are two key elements to get a nearly peak FLOPS
- the last key element: **Instruction Level Parallelism (ILP)** of superscalar processors

Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

An endeavor to nearly peak FLOPS

- measure how fast we can iterate the following loop (a similar experiment we did on GPU)

```
1  floatv a, x, c;  
2  for (i = 0; i < n; i++) {  
3      x = a * x + c;  
4  }
```

- the code performs $L \times n$ FMAs and almost nothing else (L = the number of lanes in a single SIMD variable)

Assembly

.LBB3_8:

```
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
vmadd213pd %zmm1, %zmm0, %zmm2  
addq $-8, %rax  
jne .LBB3_8
```

- the loop is unrolled eight times
- why does it take > 3 cycles to do a single `fmadd`?



Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

Latency and throughput

- our core (Ice Lake) can execute *two fmadd instructions every cycle*
- but it does *not* mean the result of `vmadd` at a line below is available in the next cycle for `vmadd` at the next line

```
1  .LBB3_8:
2    vmadd213pd %zmm1, %zmm0, %zmm2
3    vmadd213pd %zmm1, %zmm0, %zmm2
4    vmadd213pd %zmm1, %zmm0, %zmm2
5    vmadd213pd %zmm1, %zmm0, %zmm2
6    vmadd213pd %zmm1, %zmm0, %zmm2
7    vmadd213pd %zmm1, %zmm0, %zmm2
8    vmadd213pd %zmm1, %zmm0, %zmm2
9    vmadd213pd %zmm1, %zmm0, %zmm2
10   addq $-8, %rax
11   jne .LBB3_8
```

Latency and throughput

- our core (Ice Lake) can execute *two vmadd instructions every cycle*
- but it does *not* mean the result of vmadd at a line below is available in the next cycle for vmadd at the next line

```
1  .LBB3_8:  
2    vmadd213pd %zmm1, %zmm0, %zmm2  
3    vmadd213pd %zmm1, %zmm0, %zmm2  
4    vmadd213pd %zmm1, %zmm0, %zmm2  
5    vmadd213pd %zmm1, %zmm0, %zmm2  
6    vmadd213pd %zmm1, %zmm0, %zmm2  
7    vmadd213pd %zmm1, %zmm0, %zmm2  
8    vmadd213pd %zmm1, %zmm0, %zmm2  
9    vmadd213pd %zmm1, %zmm0, %zmm2  
10   addq $-8, %rax  
11   jne .LBB3_8
```

- *what you need to know:*
 - “two vmadd instructions every cycle” refers to the *throughput*
 - each instruction has a specific *latency* (≥ 1 cycle)

Latencies/throughput

instruction	Haswell	Broadwell	Skylake
fp add	3	3	4/2
fp mul	5	3	4/2
fp fmadd	5	5	4/2
typical integer ops	1	1	1/> 2
...

Valuable resources for detailed analyses

- Software optimization resources by Agner
 - *The microarchitecture of Intel, AMD and VIA CPUs: An optimization guide for assembly programmers and compiler makers*
 - *Instruction tables: Lists of instruction latencies, throughputs and micro-operation breakdowns for Intel, AMD and VIA CPUs*
- Intel Intrinsics Guide
- Intel Architecture Code Analyzer (later)

Our code in light of latencies


- in our code, a `vmadd` uses the result of the immediately preceding `vmadd`
- there are *dependencies* between them
- *that was obvious from the source code too*

```
1 .LBB3_8:  
2   vmadd213pd %zmm1, %zmm0, %zmm2  
3   vmadd213pd %zmm1, %zmm0, %zmm2  
4   ...  
5   vmadd213pd %zmm1, %zmm0, %zmm2  
6   vmadd213pd %zmm1, %zmm0, %zmm2  
7   addq $-8, %rax  
8   jne .LBB3_8
```

```
1   for (i = 0; i < n; i++) {  
2       x = a * x + c;  
3   }
```


Conclusion:

the loop can't run faster than 4 cycles/iteration



CPU clocks vs. reference clocks

- CPU changes clock frequency depending on the load (DVFS)
- reference clock runs at the same frequency (it is always proportional to the absolute time)
- an instruction takes a specified number of *CPU clocks*, not reference clocks
- the CPU clock is more predictable and thus more convenient for a precise reasoning of the code



Contents


- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

How to overcome latencies?

- increase parallelism (no other ways)!


How to overcome latencies?

- increase parallelism (no other ways)!
- you *can't* make a serial chain of dependent computation run faster than determined by latencies



How to overcome latencies?

- increase parallelism (no other ways)!
- you *can't* make a serial chain of dependent computation run faster than determined by latencies




- you *can* only increase *throughput*, by running multiple *independent* chains



How to overcome latencies?

- increase parallelism (no other ways)!
- you *can't* make a serial chain of dependent computation run faster than determined by latencies



- you *can* only increase *throughput*, by running multiple *independent* chains




- we expect the following to finish in the same number of cycles as the original one, despite it performs twice as many flops

```
1  for (i = 0; i < n; i++) {  
2      x0 = a * x0 + c;  
3      x1 = a * x1 + c;  
4  }
```

Increase the number of chains further ...

- we expect to reach peak FLOPS with $\geq 2/(1/4) = 8$ chains (i.e., $nv \geq 8$)


```
1 long axpy_simd_c( ... ) {  
2   for (long i = 0; i < n; i++) {  
3     for (long j = 0; j < nv; j++) {  
4       X[j] = a * X[j] + c;  
5     } } }
```



- note: the above reasoning assumes a compiler's smartness
- in particular, $X[j] = a * X[j] + c$ is compiled into an FMA instruction on registers without load/store instructions (i.e., each of $X[0]$, ..., $X[7]$ gets assigned a register)

Results

a compile-time constant number of variables



```
1 for (i = 0; i < n; i++) {  
2     x0 = a * x0 + b;  
3     x1 = a * x1 + b;  
4     ...  
5 }
```

chains	clocks/iter	flops/clock
1	4.010	7.979
2	4.003	15.987
3	4.013	23.916
4	4.043	31.653
5	4.043	39.568
6	4.047	47.439
7	4.157	53.878
8	5.044	50.751
9	4.621	62.314
10	5.057	63.270
11	5.549	63.427
12	6.076	63.194
13	6.573	63.283
14	7.022	63.794
15	7.552	63.558

Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

Superscalar processors

how modern aggressive superscalar processors work:

Superscalar processors

how modern aggressive superscalar processors work:

- instruction decoding goes much ahead of actual executions

Superscalar processors

how modern aggressive superscalar processors work:

- instruction decoding goes much ahead of actual executions
- the actual execution of an instruction does not happen until, and happens as soon as, *its operands and execution resources are ready* (out of order execution)

Superscalar processors

how modern aggressive superscalar processors work:

- instruction decoding goes much ahead of actual executions
- the actual execution of an instruction does not happen until, and happens as soon as, *its operands and execution resources are ready* (out of order execution)
- \Rightarrow as a crude approximation, performance is constrained by

Superscalar processors

how modern aggressive superscalar processors work:

- instruction decoding goes much ahead of actual executions
- the actual execution of an instruction does not happen until, and happens as soon as, *its operands and execution resources are ready* (*out of order execution*)
- \Rightarrow as a crude approximation, performance is constrained by
 - *latency*: imposed by *dependencies* between instructions

Superscalar processors


how modern aggressive superscalar processors work:

- instruction decoding goes much ahead of actual executions
- the actual execution of an instruction does not happen until, and happens as soon as, *its operands and execution resources are ready* (out of order execution)
- \Rightarrow as a crude approximation, performance is constrained by
 - *latency*: imposed by *dependencies* between instructions
 - *throughput*: imposed by execution resources of the processor (e.g., two fmadds/cycle)

A general theory of workload performance on aggressive superscalar machines

- *dependency* constrains how fast a computation can proceed, even if there are infinite number of execution resources
- increase the number of independent computations and you increase *throughput*, until it hits the limit of execution resources

a compile-time constant number of variables



A more general understanding about *throughput* limits

- *what you need to know:*

all instructions have their own throughput limits (just like FMA), due to execution resources

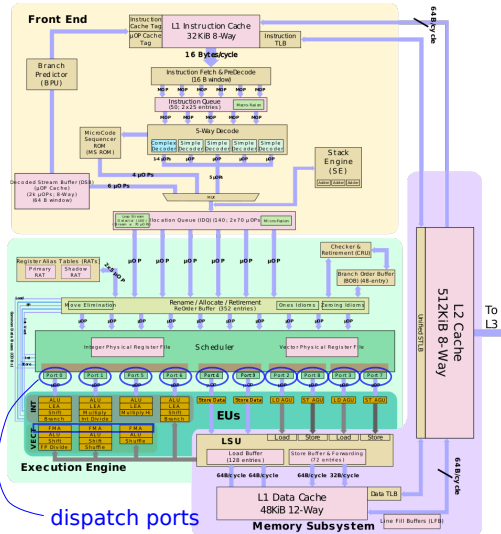
- some examples of recent Intel CPUs

instruction	Broadwell	Skylake SP	Ice Lake SP
fp add/mul/fmadd	2	2	2
load	2	2	2
store	1	1	2
typical integer ops	4	4	4
...

- e.g., a loop containing 10 load instructions takes $\geq 10/2 = 5$ cycles/iteration
- different but similar instructions may use the same execution resource so may be subject of the same limitation
- a more general reasoning \Rightarrow *dispatch ports*

Dispatch ports

- each instruction (μ -operation) is dispatched to a specific execution unit through a *dispatch port*
- each port can take only a single operation per cycle
- this determines the throughput of all instructions that go to that port
- *with destination ports of instructions, one can calculate the throughput limit of a given loop*



Chipwika - Sunny cove architecture, CC BY-SA 4.0,

<https://commons.wikimedia.org/w/index.php?curid=122557706>

LLVM Machine Code Analyzer (`llvm-mca`)

- a great tool to analyze the throughput (and latency to some extent) limit
- given a code sequence, it shows
 - latency and
 - dispatch portof each instruction and, based on them calculates the number of cycles per iteration,
- under some simplifying assumptions
 - the given sequence repeats many times
 - no cache misses (!)
 - no dependencies through memory (load does not depend on earlier stores)
 - no branch misprediction
- \Rightarrow a great tool to analyze the innermost, straight sequence of instructions without branches (basic blocks)

How to use `llvm-mca`

- 1 generate assembly (get `program.s`) by, e.g.,

```
1 clang -O3 -mavx512f -mfma ... program.c -S
```

- 2 find the loop you want to analyze in the assembly
- 3 sandwich it by `# LLVM-MCA-BEGIN` and `# LLVM-MCA-END`

```
1 # LLVM-MCA-BEGIN
2 .L123
3     ...
4     ...
5     jne .L123
6 # LLVM-MCA-END
```

- 4 run `llvm-mca` tool on the assembly code

```
1 llvm-mca program.s
```

How to use `llvm-mca`

- it shows
 - latency of each instruction
 - dispatch port used by each instruction

and how many instructions use each of the dispatch ports
(therefore the throughput limit of the loop)

- with `--timeline` option,

```
1  llvm-mca --timeline program.s
```

it also shows when each instruction gets decoded, dispatched,
and finished (particularly instructive)

Example

- input (assembly)

```
1  # LLVM-MCA-BEGIN
2  .LBB3_8:
3  # xmm0 = (xmm1 * xmm0) + xmm2
4  vfmadd213sd %xmm2, %xmm1, %xmm0
5  vfmadd213sd %xmm2, %xmm1, %xmm0
6  vfmadd213sd %xmm2, %xmm1, %xmm0
7  vfmadd213sd %xmm2, %xmm1, %xmm0
8  vfmadd213sd %xmm2, %xmm1, %xmm0
9  vfmadd213sd %xmm2, %xmm1, %xmm0
10 vfmadd213sd %xmm2, %xmm1, %xmm0
11 vfmadd213sd %xmm2, %xmm1, %xmm0
12 addq $-8, %rax
13 jne .LBB3_8
14 # LLVM-MCA-END
```

Example

- output (dispatch port used by each instruction)

```
1 Resource pressure by instruction:
2 [0] [1] [2] [3] [4] .. [11] Instructions:
3 - - 0.99 0.01 - - vfmadd213sd %xmm2, %xmm1, %xmm0
4 - - - 1.00 - - vfmadd213sd %xmm2, %xmm1, %xmm0
5 - - 0.99 0.01 - - vfmadd213sd %xmm2, %xmm1, %xmm0
6 - - - 1.00 - - vfmadd213sd %xmm2, %xmm1, %xmm0
7 - - 1.00 - - - vfmadd213sd %xmm2, %xmm1, %xmm0
8 - - - 1.00 - - vfmadd213sd %xmm2, %xmm1, %xmm0
9 - - 1.00 - - - vfmadd213sd %xmm2, %xmm1, %xmm0
10 - - - 1.00 - - vfmadd213sd %xmm2, %xmm1, %xmm0
11 - - - 0.01 - - addq $-8, %rax
12 - - 0.04 - - - jne .LBB3_8
```

Example

- output (timeline)

```
1      ...
2  D=====...==eEEEER . . . vfmadd213sd %xmm2, %xmm1, %xmm0
3  .D=====...=====eEEEER . . . vfmadd213sd %xmm2, %xmm1, %xmm0
4  .D=====...=====eEEEER. . . vfmadd213sd %xmm2, %xmm1, %xmm0
5  .DeE--...-----R. . . addq $-8, %rax
6  .D=eE-...-----R. . . jne .LBB3_8
7  .D=====...=====eEEEER . . . vfmadd213sd %xmm2, %xmm1, %xmm0
8  .D=====...=====eEEEER . vfmadd213sd %xmm2, %xmm1, %xmm0
9  . D==...=====eEEEER vfmadd213sd %xmm2, %xmm1, %xmm0
10     ...
```


Contents

- 1 Introduction
- 2 An endeavor to nearly peak FLOPS
- 3 Latency limit
- 4 Overcoming latency
- 5 Superscalar processors
- 6 A simple yet fairly fast single-core matrix multiply

Developing near peak FLOPS matrix multiply

- let's develop a (single core) matrix multiply that runs at fairly good FLOPS on Ice Lake
- it is a great application of the concept you have just learned


$$C = A * B + C$$



Developing near peak FLOPS matrix multiply

- let's develop a (single core) matrix multiply that runs at fairly good FLOPS on Ice Lake
- it is a great application of the concept you have just learned


$$C = A * B + C$$



A few convenient assumptions

- we add assumptions that M , N , and K are multiple of certain numbers along the way, (don't worry about "remainder" rows/columns)
- we assume matrix sizes are conveniently small (don't worry about memory access cost, which is actually a significant factor to design matrix multiply for larger matrices)
- multiplication of larger (and unknown size) matrices can be built on top of this

Step 1: Baseline code



```
1 $ mm_base 8 32 192
2 M = 8, N = 32, K = 192
3 L : 16
4 A : 8 x 192 (ld=192) 6144 bytes
5 B : 192 x 32 (ld=32) 24576 bytes
6 C : 8 x 32 (ld=32) 1024 bytes
7 total = 31744 bytes
8 repeat : 20346 times
9 perform 1000046592 fmas ... done
10 2844287815 clocks
11 0.351598 fmas/cycle
```


```
1 for (i = 0; i < M; i++)
2   for (j = 0; j < N; j++)
3     for (k = 0; k < K; k++)
4       C(i,j) += A(i,k) * B(k,j);
```

- it runs at ≈ 2.8 clocks / innermost loop

Step 1: analysis

- latency limit : latency of FMA
 - the reason why it's slightly *smaller* than 4 is there are some overlaps between different elements of C
 - if you set $M = N = 1$ and K large, it's almost exactly 4
- throughput limit : not important
- achieved performance : $1000046592 \text{ fmas} / 2844287815 \text{ cycles}$
 $\approx 0.4 \text{ fmas/cycle}$

Step 2: Vectorization




```
1 $ mm_simd 8 32 192
2 M = 8, N = 32, K = 192
3 L : 16
4 A : 8 x 192 (ld=192) 6144 bytes
5 B : 192 x 32 (ld=32) 24576 bytes
6 C : 8 x 32 (ld=32) 1024 bytes
7 total = 31744 bytes
8 repeat : 20346 times
9 perform 1000046592 fmas ... done
10 180175475 clocks
11 5.550404 fmas/cycle
```

```
1 for (i = 0; i < M; i++)
2   for (j = 0; j < N; j += L)
3     for (k = 0; k < K; k++)
4       C(i,j:j+L) += A(i,k) * B(k,j:j+L);
```

- assumption: N is a multiple of SIMD lanes (L)
- it still runs at ≈ 2.8 clocks / innermost iteration

Step 2: Vectorization



```
1 $ mm_simd 8 32 192
2 M = 8, N = 32, K = 192
3 L : 16
4 A : 8 x 192 (ld=192) 6144 bytes
5 B : 192 x 32 (ld=32) 24576 bytes
6 C : 8 x 32 (ld=32) 1024 bytes
7 total = 31744 bytes
8 repeat : 20346 times
9 perform 1000046592 fmas ... done
10 180175475 clocks
11 5.550404 fmas/cycle
```

```
1 for (i = 0; i < M; i++)
2   for (j = 0; j < N; j += L)
3     for (k = 0; k < K; k++)
4       C(i,j:j+L) += A(i,k) * B(k,j:j+L);
```


- assumption: N is a multiple of SIMD lanes (L)
- it still runs at ≈ 2.8 clocks / innermost iteration

Step 2: analysis

- the speed is still limited by latency
- the only difference is that each iteration now performs 16 fmas (as opposed to an fma)
- achieved throughput :

$$1000046592 \text{ fmas} / 180175475 \text{ cycles} \approx 5.5 \text{ fmas/cycle}$$

Step 3: increase parallelism!




```
1 $ ./mm_simd_ilp 8 32 192
2 M = 8, N = 32, K = 192
3 L : 16
4 A : 8 x 192 (ld=192) 6144 bytes
5 B : 192 x 32 (ld=32) 24576 bytes
6 C : 8 x 32 (ld=32) 1024 bytes
7 total = 31744 bytes
8 repeat : 20346 times
9 perform 1000046592 fmas ... done
10 64836630 clocks
11 15.424099 fmas/cycle
```

- update bM vector elements of C concurrently

```
1 for (i = 0; i < M; i += bM)
2   for (j = 0; j < N; j += L)
3     for (k = 0; k < K; k++)
4       for (di = 0; di < bM; di++)
5         C(i+di,j:j+L) += A(i+di,k) * B(k,j:j+L);
```

- Ice Lake requires $bM \geq 8$ to reach peak FLOPS

Step 3: increase parallelism!




```
1 $ ./mm_simd_ilp 8 32 192
2 M = 8, N = 32, K = 192
3 L : 16
4 A : 8 x 192 (ld=192) 6144 bytes
5 B : 192 x 32 (ld=32) 24576 bytes
6 C : 8 x 32 (ld=32) 1024 bytes
7 total = 31744 bytes
8 repeat : 20346 times
9 perform 1000046592 fmas ... done
10 64836630 clocks
11 15.424099 fmas/cycle
```

- update bM vector elements of C concurrently

```
1 for (i = 0; i < M; i += bM)
2   for (j = 0; j < N; j += L)
3     for (k = 0; k < K; k++)
4       for (di = 0; di < bM; di++)
5         C(i+di,j:j+L) += A(i+di,k) * B(k,j:j+L);
```

- Ice Lake requires $bM \geq 8$ to reach peak FLOPS

Step 3: analysis




```
1 for (i = 0; i < M; i += bM)
2   for (j = 0; j < N; j += L)
3     for (k = 0; k < K; k++)
4       for (di = 0; di < bM; di++)
5         C(i+di,j:j+L) += A(i+di,k) * B(k,j:j+L);
```

- the for loop at line 4 performs
 - bM loads (broadcasts) for $A(i+di,k)$
 - 1 load for $B(k,j:j+L)$
 - bM FMAs
- the load/broadcast throughput = 2 per cycle
- to achieve 2 FMAs/cycle, we must have

the number of broadcast \leq the number of FMAs

Step 4: Reuse an element of A



```

1  $ mm_simd_lip_4x2 8 32 192
2  M = 8, N = 32, K = 192
3  L : 16
4  A : 8 x 192 (ld=192) 6144 bytes
5  B : 192 x 32 (ld=32) 24576 bytes
6  C : 8 x 32 (ld=32) 1024 bytes
7  total = 31744 bytes
8  repeat : 20346 times
9  perform 1000046592 fmas ... done
10 38635137 clocks
11 25.884381 fmas/cycle
12

```

- update $bM' \times bN$ block rather than $bM \times 1$

```

1  for (i = 0; i < M; i += bM')
2    for (j = 0; j < N; j += bN * L)
3      for (k = 0; k < K; k++)
4        for (di = 0; di < bM'; di++)
5          for (dj = 0; dj < bN * L; dj += L)
6            C(i+di,j+dj:j+dj+L) += A(i+di,k) * B(k,j+dj:j+L);

```

Step 4: Analysis


- the for loop at line 4 performs
 - bM' loads (broadcast) for $A(i+di, k)$
 - bN loads for $B(k, j:j+L)$
 - $bM' \times bN$ SIMD FMAs
- the minimum requirement for it to achieve the peak FLOPS is $bM' \times bN \geq 8$
- in the experiments, when we set $bM' = 8$ and $bN = 2$, it gets 25 fmas/cycle ($\approx 80\%$ of the peak)
- we need to note that this happens only when the matrix is small ($M = 8, N = 32, K = 192$) and we repeat it many times
- the issue for large matrices will be the next topic

Takeaways (1)


- peak FLOPS of many recent Intel CPUs = “execute two `fmadds` every cycle” (*no other combinations*)
 - other processors have different limits, but the basics is the same
 - cf. NVIDIA GPUs = “execute two warps (each doing `fmadd`) every cycle”
- single-core performance is not about reducing the number of instructions
- it's about how to increase parallelism
 - CPU : $\text{SIMD} \times \text{ILP}$
 - GPU : threads, threads, threads, ...
 - but the internal machinery is similar (warp \approx SIMD, ILP \sim warps in an SM)
 - how they expose parallelism to the programmer is different

Takeaways (2)

- dependent instructions incur latencies and hinder parallelism



- independent instructions are executed in parallel, up to throughput limits



- throughput limits are determined by dispatch ports