SIMD Programming

Kenjiro Taura

Contents

Contents

SIMD : basic concepts

- SIMD : single instruction multiple data
- a *SIMD register* (or a *vector register*) can hold many values (2 - 16 values or more) of a single type
- a *SIMD instruction* is an instruction that can apply (typically the same) operation on all or some values on a SIMD register(s)
- each value in a SIMD register is called a *SIMD lane* or simply a *lane*
- they are indispensable tools for CPUs to get performance

A SIMD register

...

lane

Evolving Intel instruction set

- Recent processors increasingly rely on SIMD as an energy efficient way to boost peak FLOPS

Microarchitecture	ISA	throughput (per clock)	vector width (SP)	max SP flops/cycle /core
Nehalem	SSE	1 add + 1 mul	4	8
Sandy Bridge	AVX	1 add + 1 mul	8	16
Haswell	AVX2	2 fmas	8	32
Ice Lake	AVX-512	2 fmas	16	64

- ISA : Instruction Set Architecture
- vector width : the number of single precision (SP) operands
- fma : fused multiply-add instruction
- e.g., Peak FLOPS of a machine having $2 \times$ Intel Xeon Gold 6130 (2.10GHz, 32 cores) = 8.6 TFLOPS
- no SIMD? $\rightarrow$ can tap *at most 1/16* of SP peak performance on machines having AVX-512

Intel SIMD instructions at a glance

Some example AVX-512F (a subset of AVX-512) instructions

operation	syntax	C-like expression
multiply	`vmulps %zmm0,%zmm1,%zmm2`	`zmm2 = zmm1 * zmm0`
add	`vaddps %zmm0,%zmm1,%zmm2`	`zmm2 = zmm1 + zmm0`
fmadd	`vfmadd132ps %zmm0,%zmm1,%zmm2`	`zmm2 = zmm0*zmm2+zmm1`
load	`vmovups 256(%rax),%zmm0`	`zmm0 = *(rax+256)`
store	`vmovups %zmm0,256(%rax)`	`*(rax+256) = zmm0`

- `zmm0` ... `zmm31` are 512 bit registers; each can hold
  - 16 single-precision (`float` of C; 32 bits) or
  - 8 double-precision (`double` of C; 64 bits)
    floating point numbers
- `XXXps` stands for *packed single precision*

xmm, ymm and zmm registers

- ISA and available registers

ISA	registers
SSE	`xmm0`, `...xmm15`
AVX	`{x,y}mm0,...{x,y}mm15`
AVX-512	`{x,y,z}mm0,...{x,y,z}mm31`

- registers and their widths *(vector widths)*

register names	register width (bits)
xmm$i$	128
ymm$i$	256
zmm$i$	512

- xmm$i$, ymm$i$ and zmm$i$ are *aliased*

Intel SIMD instructions at a glance

- look at register names (x/y/z) and the last two characters of a mnemonic (p/s and s/d) to know what an instruction operates on

	operands	vector /scalar?	ISA
`vmulss %xmm0,%xmm1,%xmm2`	1 SPs	scalar	SSE
`vmulsd %xmm0,%xmm1,%xmm2`	1 DPs	scalar	SSE
`vmulps %xmm0,%xmm1,%xmm2`	4 SPs	vector	SSE
`vmulpd %xmm0,%xmm1,%xmm2`	2 DPs	vector	SSE
`vmulps %ymm0,%ymm1,%ymm2`	8 SPs	vector	AVX
`vmulpd %ymm0,%ymm1,%ymm2`	4 DPs	vector	AVX
`vmulps %zmm0,%zmm1,%zmm2`	16 SPs	vector	AVX-512
`vmulpd %zmm0,%zmm1,%zmm2`	8 DPs	vector	AVX-512

- . . . ss : *s*calar *single* precision
- . . . sd : *s*calar *double* precision
- . . . ps : *p*acked *single* precision
- . . . pd : *p*acked *double* precision

Applications/limitations of SIMD

- SIMD is good at parallelizing computations doing *almost exactly* the same series of instructions on contiguous data
- $\Rightarrow$ generally, main targets are simple loops whose index values can be easily identified

```
1  for (i = 0; i < n; i++) {
2    S(i);
3  }
```

$\Rightarrow$

```
1  for (i = 0; i + L <= n; i += L) {
2    S(i : i + L);
3  }
4  for (; i < n; i++) { /* remainder iterations */
5    S(i);
6  }
```

L is the SIMD width

Contents

Several ways to use SIMD

- auto vectorization
  - loop vectorization
  - basic block vectorization
- language extensions/directives for SIMD
  - SIMD directives for loops (OpenMP 4.0/OpenACC)
  - SIMD-enabled functions (OpenMP 4.0/OpenACC)
  - array languages (Cilk Plus)
  - specially designed languages
- vector types
  - GCC vector extensions
  - Boost.SIMD
- intrinsics
- assembly programming

Contents

Auto loop vectorization

- write scalar loops and hope the compiler does the job
- e.g.,

```
1  void axpy_auto(float a, float * x, float c, long m) {
2    for (long j = 0; j < m; j++) {
3      x[j] = a * x[j] + c;
4    }
5  }
```

- compile and run

```
1  $ clang -o simd_auto -mavx512f -mfma -O3 simd_auto.c
```

- `-mavx512f` `-mfma` say "should use AVX-512F and FMA instructions" (better to be explicit for the time being)
- `-O3` increases the optimization level (so the compiler should work hard to vectorize it)
- read the notebook about options of other compilers (NVIDIA and GCC)

How to know if the compiler vectorized it?

- there are options useful to know whether a loop is successfully vectorized and if not, why not

	report options
Clang	`-R{pass,pass-missed}=loop-vectorize`
NVIDIA	`-M{info,neginfo}=vect`
GCC	`-fopt-info-vec-{optimized,missed}`

- *but don't hesitate to dive into assembly code*
  - make `-S` option your friend
  - a trick: *enclose loops with inline assembler comments to easily locate assembly code for the loop*

```
1  asm volatile ("# xxxxxx loop begins");
2  for (i = 0; i < n; i++) {
3    ... /* hope to be vectorized */
4  }
5  asm volatile ("# xxxxxx loop ends");
```

Contents

OpenMP SIMD constructs

- `simd` pragma
  - directive to vectorize for loops
  - syntax restrictions similar to `omp for` pragma apply
- `declare simd` pragma
  - instructs the compiler to generate vectorized versions of a function
  - with it, loops with function calls can be vectorized

`simd` pragma

- basic syntax (similar to `omp for`):

```
#pragma omp simd clauses
for (i = ...; i < ...; i += ...)
    S
```

- clauses
    - `aligned(`*var,var,...* : *align*`)`
    - `uniform(`*var,var,...*`)` says variables are loop invariant
    - `linear(`*var,var,...* : *stride*`)` says variables have the specified stride between consecutive iterations

`simd` pragma

```
1  void axpy_omp(float a, float * x, float c, long m) {
2  #pragma omp simd
3    for (long j = 0; j < m; j++) {
4      x[j] = a * x[j] + c;
5    }
6  }
```

- note: there are no points in using `omp simd` here, when auto vectorization does the job
- in general, `omp simd` declares "you don't mind that the vectorized version is not the same as non-vectorized version"

`simd` pragma to vectorize programs explicitly

- computing an inner product:

```
1  void inner_omp(float * x, float * y, long m) {
2    float c = 0;
3  #pragma omp simd reduction(c:+)
4    for (long j = 0; j < m; j++) {
5      c += x[j] * y[j];
6    }
7  }
```

- note that the above loop is unlikely to be auto-vectorized,
  due to dependency through c

`declare simd` pragma

- when given before a function definition, vectorizes a function body
- when given before a function declaration, tells the compiler a vectorized version of the function is available
- basic syntax (similar to `omp for`):

```
1  #pragma omp declare simd clauses
2  function definition or declaration
```

- clauses
  - those for `simd` pragma
  - `notinbranch`
  - `inbranch`

Reasons that a vectorization fails

- potential aliasing makes auto vectorization difficult/impossible
- complex control flows make vectorization impossible or less profitable
- non-contiguous data accesses make vectorization impossible or less profitable

  giving hints to the compiler sometimes (not always) addresses the problem

Aliasing and auto vectorization

- "auto" vectorizer succeeds only when the compiler can guarantee a vectorized version produces an *identical result* with a non-vectorized version
- vectorization of loops operating on two or more arrays is often invalid if they point to be the same array

```
1  for (i = 0; i < m; i++) {
2    y[i] = a * x[i] + c;
3  }
```

what if, say, &y[i] = &x[i+1]?

  - N.B., good compilers generate code that first checks x[i:i+*L*] and y[i:i+*L*] overlap
- if you know they don't overlap, you can make that explicit
- `restrict` keyword, introduced by C99, does just that

`restrict` keyword

- annotate parameters of pointer type with `restrict`, if you
  know they never point to the same data

```c
void axpy_auto(float a, float * restrict x, float c,
          float * restrict y, long m) {
    for (long j = 0; j < m; j++) {
        y[j] = a * x[j] + c;
    }
}
```

- you need to specify `-std=gnu99` (C99 standard)

```
$ gcc -march=native -O3 -S a.c -std=gnu99 -fopt-info-vec-optimized
  ...
a.c:5: note: LOOP VECTORIZED.
a.c:1: note: vectorized 1 loops in function.
  ...
```

Control flows within an iteration — conditionals

- a conditional execution (e.g., if statement) within an iteration requires a statement to be executed only for a part of SIMD lanes

```
1  void loop_if(float a, float * restrict x, float b,
2              float * restrict y, long n) {
3  #pragma omp simd
4    for (long i = 0; i < n; i++) {
5      if (x[i] < 0.0) {
6        y[i] = a * x[i] + b;
7      }
8    }
9  }
```

- AVX-512 supports *predicated execution (execution mask)* for that

Control flows within an iteration — nested loops

- a nested loop within an iteration causes a similar problem with conditional executions

```
1   void loop_loop(float a, float * restrict x, float b,
2                  float * restrict y, long n) {
3   #pragma omp simd
4     for (long i = 0; i < n; i++) {
5       y[i] = x[i];
6       for (long j = 0; j < end; j++) {
7         y[i] = a * y[i] + b;
8       }
9     }
10  }
```

- if *end* depends on i (SIMD lanes), it requires a predicated execution

Control flows within an iteration — function calls

- if an iteration has an unknown (not inlined) function call, almost no chance that the loop can be vectorized
  - the function body would have to be executed by scalar instructions anyways

```
void loop_fun(float a, float * restrict x, float b,
              float * restrict y, long n) {
#pragma omp simd
  for (long i = 0; i < n; i++) {
    f(a, x, b, y, i);
  }
}
```

- you can declare that f has a vectorized version with #pragma omp declare simd (with such a definition, of course)

```
#pragma omp declare simd uniform(a, x, b, y) linear(i:1) notinbranch
void f(float a, float * restrict x, float b, float * restrict y, long i);
```

Non-contiguous data accesses

- ordinary vector load/store instructions access a contiguous addresses

```
1  vmovups (a),%zmm0
```

loads zmm0 with the contiguous 64 bytes from address a

- $\rightarrow$ they can be used only when iterations next to each other access addresses next to each other

Non-contiguous data accesses

- that is, they cannot be used for

```
1  void loop_stride(float a, float * restrict x, float b,
2                   float * restrict y, long n) {
3  #pragma omp simd
4    for (long i = 0; i < n; i++) {
5      y[i] = a * x[2 * i] + b;
6    }
7  }
```

let alone

```
1  void loop_random(float a, float * restrict x, float b,
2                   float * restrict y, long n) {
3  #pragma omp simd
4    for (long i = 0; i < n; i++) {
5      y[i] = a * x[i * i] + b;  // or x[idx[i]]
6    }
7  }
```

- AVX-512 supports *gather* instructions for such data accesses

Non-contiguous stores

- what about store

```
1  void loop_random_store(float a, float * restrict x, long * idx, float b,
2                         float * restrict y, long n) {
3  #pragma omp simd
4    for (long i = 0; i < n; i++) {
5      y[idx[i]] += a * x[i] + b;
6    }
7  }
```

- AVX-512 supports *scatter* instructions for such data accesses
- it is your responsibility to guarantee `idx[i:i+L]` do not point to the same element

High level vectorization: summary and takeaway

- CPUs (especially recent ones) have necessary tools
  - arithmetic $\rightarrow$ vector arithmetic instructions
  - load $\rightarrow$ vector load and gather instructions
  - store $\rightarrow$ vector store and scatter instructions
  - if and loops $\rightarrow$ predicated executions
- generally, the compiler is behind CPUs; whether the compiler is able to use them is another story
- become a friend of compiler reports and assembly (`-S`)

Contents

Vector types

- many compilers extend C by allowing you to define a type that explicitly represents a vector of values

```
1  typedef float floatv __attribute__((vector_size(64)));
```

- you can use familiar arithmetic expressoins on vector types

```
1  floatv x, y, z;
2  z += x * y;
```

- Clang/NVIDIA/GCC allow you to mix scalars and vectors

```
1  float a, b;
2  floatv x, y;
3  y = a * x + b;
```

- you can combine them with *intrinsics* I'll get to later
- for reasons I don't get into, a better definition is

```
1  typedef float floatv __attribute__((vector_size(64),
2                                       __may_alias__,
3                                       aligned(sizeof(float))));
```

An example using vector extension

- scalar code

```
1   for (long i = 0; i < n; i++) {
2       y[i] = a * x[i] + b;
3   }
```

- pseudo code (assume L | n (L divides n))

```
1   for (long i = 0; i < n; i += L) {
2       y[i:i+L] = a * x[i:i+L] + b;
3   }
```

- a function or macro (V) implementing x[i:i+L]

```
1   /* take the address, cast it to (floatv *) and deref it */
2   #define V(lv) (*((floatv*)&(lv)))
```

- it is then

```
1   for (long i = 0; i < n; i += L) {
2       V(y[i]) = a * V(x[i]) + b;
3   }
```

Dealing with remainder iterations

- when L $\not|$ n, run remainders after the vectorized version

```
1   long i;
2   for (i = 0; i + L <= n; i += L) {
3     V(y[i]) = a * V(x[i]) + b;
4   }
5   for (    ; i < n; i++) {
6     y[i] = a * x[i] + b;
7   }
```

- manually doing this is tedious . . .
- make n a multiple of L when the problem allows it (otherwise do the tedious work)

Make a vector value from scalar value(s)

- you typically make a vector value from an array of scalars

```
1  float * a = ...;
2  floatv v = *((*floatv)&a[i]);
```

- a macro/function like the following makes the life better

```
1  floatv& V(float& lv) { return *((floatv*)(&lv)); } // C++
2  #define V(lv) (*((floatv*)&(lv)))                   // C
```

with which we can write

```
1  float * a = ...;
2  floatv v = V(a[i]);
```

... and vice versa

- you typically store a vector value to an array of scalars

```
1  float * a = ...;
2  floatv v = ...;
3  V(a[i]) = v;
```

  and get individual scalars from the array

- you can access a particular lane of a vector directly, as if a vector is a C array. e.g.,

```
1  floatv v;
2  float s = v[3];
```

- but a CPU generally lacks instructions to access a lane designated by a value not known at the compile time. e.g.,

```
1  floatv v; int i = ...;
2  float s = v[i];
```

  it might be essentially doing the former each time you access an element, so might be very inefficient

Contents

Vector intrinsics

- processor/platform-specific functions and types
- on x86 processors, put this in your code

```
1  #include <x86intrin.h>
```

and you get
- a set of available vector types
- a lot of functions operating on vector types

- bookmark "Intel Intrinsics Guide" (https://software.intel.com/sites/landingpage/IntrinsicsGuide/) when using intrinsics

Vector types + intrinsics

- vectorizing a loop is largely about converting

```
1  for (i = 0; i < n; i++) {
2    S(i);
3  }
```

$\Rightarrow$

```
1  for (i = 0; i + L <= n; i += L) {
2    S(i : i + L);
3  } // + remainder code (omitted)
```

- the combination of vector types + intrinsics gives you a powerful way to *manually* vectorize code (i.e., write $S(i : i + L)$) the compiler fails to vectorize

When you want to use manual vectorization

- whenever your compiler fails, but in general
  1. a loop containing *a branch* ⇒ predicated execution + value-blending
  2. a loop accessing an array *non-contiguously* ⇒ gather + scatter
  3. a loop containing *another loop* ⇒
     - easy if all inner loops have the same trip count
     - follow the strategy for branches (tedious)

Vector intrinsics

- vector types:
    - `__m512` (512 bit vector) $\approx$ `float` $\times$ 16
    - `__m512d` (512 bit vector) $\approx$ `double` $\times$ 8
    - `__m512i` (512 bit vector) $\approx$ `long` $\times$ 8
    - there are no `int` $\times$ 16
    - similar types for 256/128 bit values (`__m256`, `__m256d`, `__m256i`, `__m128`, `__m128d` and `__m128i`
- functions operating on vector types:
    - `__mm512_xxx` (512 bit),
    - `__mm256_xxx` (256 bit),
    - `__mm_xxx` (128 bit),
    - . . .
- each function almost directly maps to a single assembly instruction

Convenient intrinsics to make a vector value from scalar value(s)

- make a uniform vector

```
1  floatv v = _mm512_set1_ps(f); // { f, f, ..., f }
```

- make an arbitrary vector

```
1  floatv v = _mm512_set_ps(f0, f1, f2, ..., f15);
```

Contents

Contents

Predicated instructions

- SIMD instructions that take a vector of boolean values (*mask*) that specifies lanes for which the instruction is executed
- results on other lanes are taken from another SIMD register (or set zero)
- e.g., an ordinary SIMD add instruction (intrinsics)
  - `__m512 __mm512_add_ps(__m512 $a$, __m512 $b$)`
    $\equiv [\; a[i] + b[i] \mid i \in 0..L \;]$
- predicated versions
  - `__m512 __mm512_maskz_add_ps(__mmask16 $k$, $a$, $b$)`
    $\equiv [\; (k[i] \; ? \; a[i] + b[i] \; : \; 0) \mid i \in 0..L \;]$
  - `__m512 __mm512_mask_add_ps(__m512 $c$, $k$, $a$, $b$)`
    $\equiv [\; (k[i] \; ? \; a[i] + b[i] \; : \; c[i]) \mid i \in 0..L \;]$

Generating a mask

- compare all values of two vectors (with $<$)
  `__mmask16` k = `_mm512_cmp_ps_mask`(a, b, `_CMP_LT_OS`)
  $\equiv [\, u[i] < v[i] \mid i \in 0..L \,]$
- you get a 16 bit *mask* that can be used for predicated execution
- search intrinsics guide for symbols to compare in other ways

A template to vectorize loops containing branches

- a loop having a branch

```
1  for (i = 0; i < n; i++) {
2    if (C(i)) {
3      T(i)
4    } else {
5      E(i)
6    } }
```

- $\Rightarrow$

```
1  for (i = 0; i + L <= n; i += L) {
2    k = C(i : i + L)
3    if (any(k)) {
4      T(i : i + L) predicated on k
5    }
6    if (any(~k)) {
7      E(i : i + L) predicated on ~k
8    } }
```

- note: values used after the original if statement are made by blending results from both branches (see next slide)

Blending values

- there are instructions specifically for blending two vectors.
  e.g., `__m512 _mm512_mask_blend_ps(k, a, b)`
  $\equiv [\, (k[i] \; ? \; a[i] \; : \; b[i]) \mid i \in i..L \,]$

- recall that predicated instructions already have a provision
  for it. e.g.,
  `__m512 _mm512_mask_add_ps(__m512 c, k, a, b)` $\equiv$
  `_mm512_mask_blend_ps(k, a + b, c)`

Example

- scalar version

```
1  for (i = 0; i < n; i++) {
2    if (i % 2 == 0) {
3      y[i] = x[i] + 1;
4    } else {
5      y[i] = x[i] * 2;
6    }
7  }
```

- $\Rightarrow$ pseudo code (assume $L \mid n$)

```
1  for (i = 0; i < n; i += L) {
2    __mmask16 k = (i:i+L % 2 == 0);
3    t = x[i:i+L] + 1;
4    y[i:i+L] = blend(~k, x[i:i+L] * 2 : t);
5  }
```

Example

- $\Rightarrow$ actual code

```
for (i = 0; i < n; i += L) {
  __m512i z = _mm512_set1_epi64(0)
  __mmask16 k = _mm512_cmp_epi64_mask(linear(i) & 1L, z, _MM_CMPINT_EQ)
  __m512i t = V(x[i]) + 1;
  V(y[i]) = _mm512_mask_mul_ps(t, ~k, V(x[i]), 2);
}
```

- linear(i) is a function (not shown) to generate a vector { i, i+1, ..., i+L-1 }
- there are C++ tricks (operator overloading) that make this code less ugly

Contents

Gather

- an instruction that can get $[\,a[i_0], a[i_1], \cdots, a[i_{L-1}]\,]$ as a vector value
- `__m512 _mm512_i32gather_ps(__m512i I, void* a, int s)` takes 16 32-bit indices I and scale s. that is,
  $\equiv [\,f(a[I[i] * s]) \mid i \in 0..L\,]$ where $f(p)$ gets the value at p as a `float` value $(\equiv *((\texttt{float}*)\&p))$
- similar versions for different index/value widths
  - 64 bit indices to gather 8 double precision (64 bit) values
    `__m512d _mm512_i64gather_pd`
  - 64 bit indices to gather 8 single precision (32 bit) values
    `__m256 _mm512_i64gather_ps`
  - 32 bit indices to gather 8 double precision (64 bit) values
    `__m512d _mm512_i32gather_pd`
- there are predicated versions as well
  (`_mm512_mask_ixxgather_ps/pd`)

Scatter

- an instruction that can assignments
  $a[i_0] = x_0; a[i_1] = x_1; \cdots ; a[i_{L-1}] = x_{L-1};$
- similar name conventions to gather
  - 32 bit indices, to get 32 bit values _mm512_i32scatter_ps
  - 64 bit indices, to get 64 bit values _mm512_i64scatter_pd
  - 64 bit indices, to get 32 bit values: _mm512_i64scatter_ps
  - 32 bit indices, to get 64 bit values: _mm512_i32scatter_pd
- you guessed it. there are masked versions
  (_mm512_mask_ixxscatter_ps/pd)

Contents

Loops having another loop inside

- consider how to vectorize the *outer* loop

```
for (i = 0; i < m; i++) {
  for (j = 0; j < limit; j++) {
    B(i)
  }
}
```

- if the trip count of the inner loop is the same across lanes (i.e., *limit* does not depend on i), then there is no particular difficulty (the compiler nevertheless often fails to vectorize it)
- more difficult is when inner loop has different trip counts depending on i

Loops having another loop inside

- a general template of scalar code

```
1  for (i = 0; i < m; i++) {
2    while (C(i)) {
3      B(i)
4    }
5  }
```

- $\Rightarrow$

```
1  for (i = 0; i < m; i += L) {
2    while (any(C(i : i + L))) {
3      B(i : i + L) predicated on C
4    }
5  }
```

Vector types and intrinsics : summary

- template

```
1  for (i = 0; i < n; i++) {
2    S(i)
3  }
```

$\rightarrow$

```
1  for (i = 0; i < n; i += L) {
2    S(i : i + L)
3  }
```

- convert every expression into its *vector* version, which contains what the original expression would have for the L consecutive iterations
- use masks to handle conditional execution and nested loops with variable trip counts
- vectorizing SpMV is challenging but possible with this approach